

Foundations of High Performance Computing

Lecture 8: MPI libraries on ORFEO and their usage



“Foundation of HPC” course

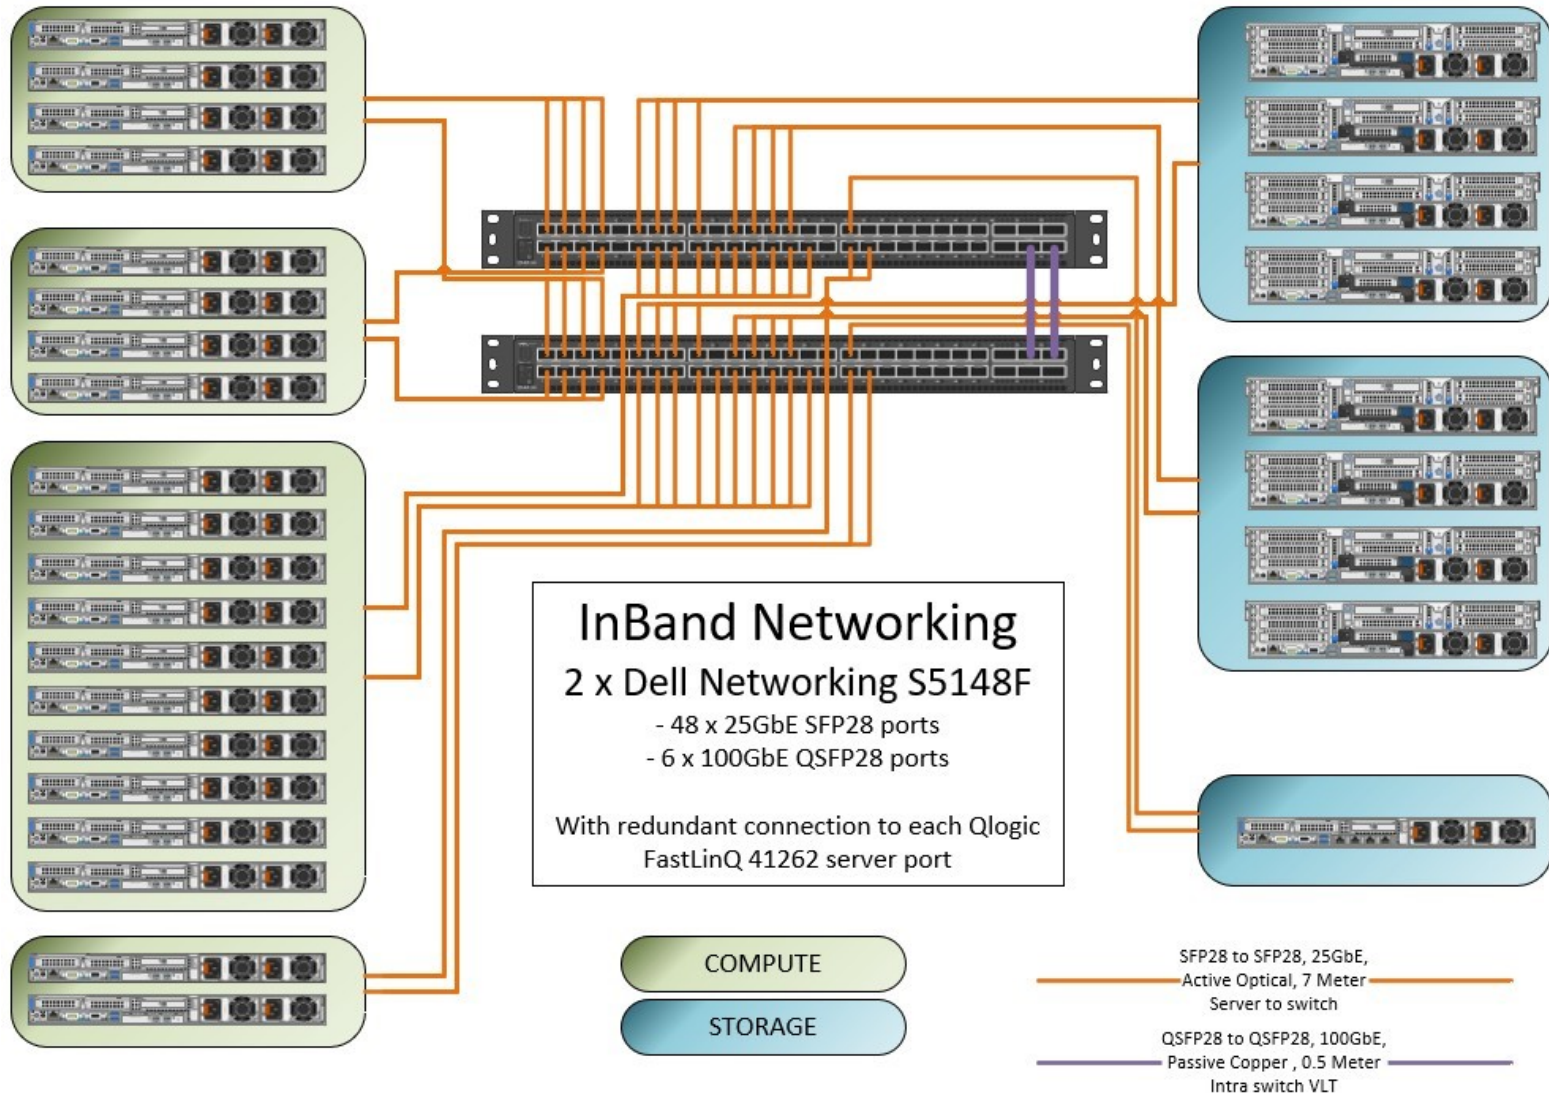
DATA SCIENCE &
SCIENTIFIC COMPUTING

2022-2023 Stefano Cozzini

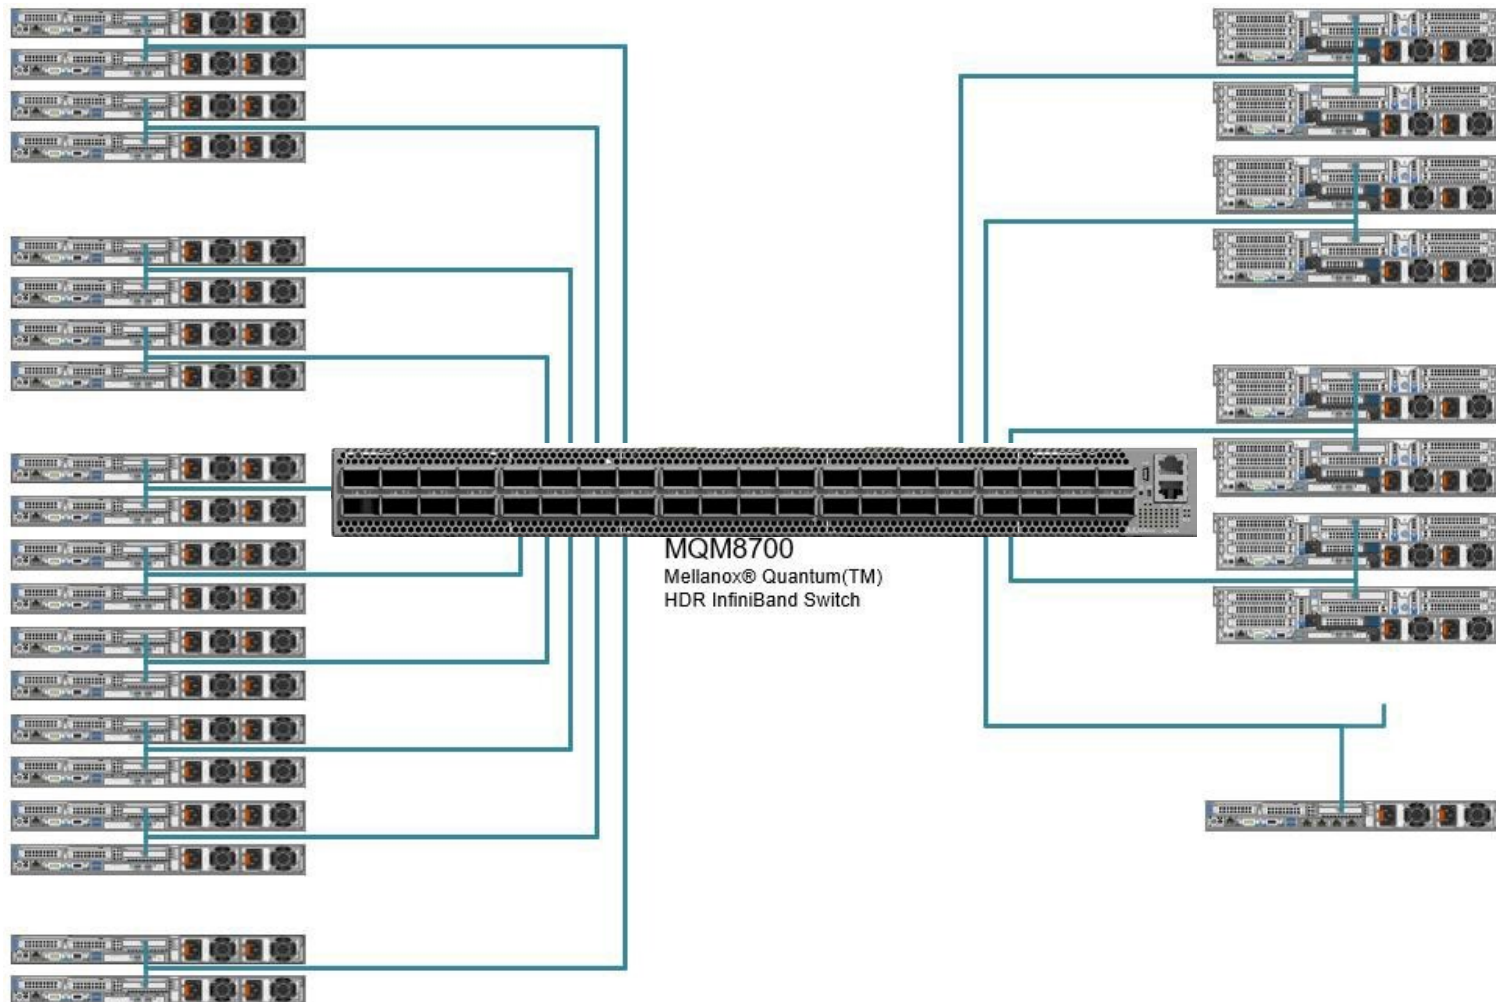
Agenda

- Recap: ORFEO networks
- Communication protocols
- MPI libraries available on ORFEO
- Measuring and understanding performance

Orfeo in band management network: 25 Gbit ethernet




Orfeo High Speed network: 100 Gbit Infiniband



IB HDR 200Gb/s to 2x100Gb/s
QSFP56 to 2xQSFP56, LSZH

ORFEO IB network




Performance

- 40 x HDR 200Gb/s ports in a 1U switch
- 80 x HDR100 100Gb/s ports (using splitter cables)
- 16Tb/s aggregate switch throughput
- Sub-130ns switch latency

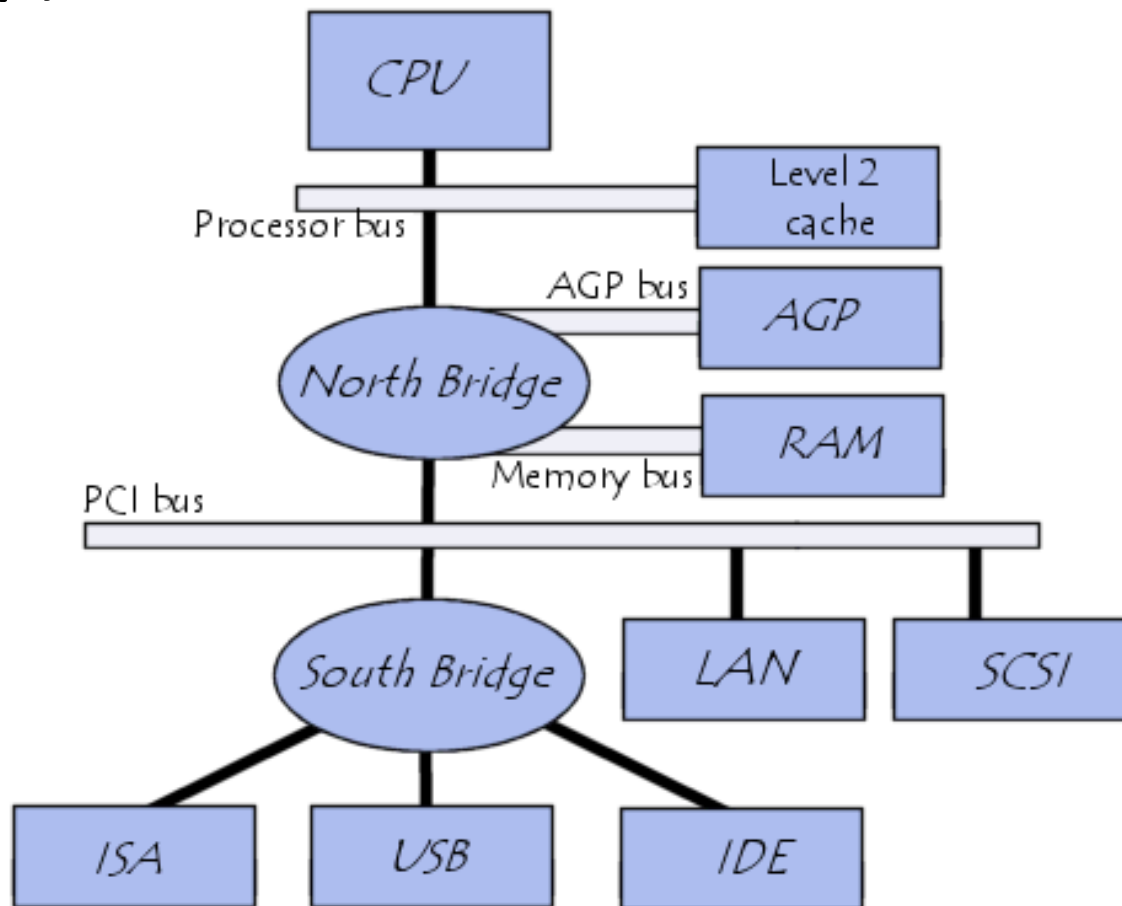


Where are the cards on the server ?

- Plugged in the PCI-e slots..



Buses within a computer (old way)



Buses on modern HPC nodes

- Peripheral Component Interconnect (PCI) buses:
 - PCI: Developed by Intel in 1992
 - several version : v3.0 last one in 2004
 - PCI-X: designed in 1999
 - 66 MHz (can be found on older servers)
 - 133 MHz (most common on modern servers)
- PCIe: designed adopted in 2004
 - version v4.0 recently released
 - Version 2.0/version 3.0 adopted on modern HPC nodes
 - Several of them on one node with different characteristics

PCI-express speed (from wikipedia)

PCI Express link performance^{[30][31]}

PCI Express version	Introduced	Line code	Transfer rate ^[i]	Throughput ^[i]				
				x1	x2	x4	x8	x16
1.0	2003	8b/10b	2.5 GT/s	250 MB/s	0.50 GB/s	1.0 GB/s	2.0 GB/s	4.0 GB/s
2.0	2007	8b/10b	5.0 GT/s	500 MB/s	1.0 GB/s	2.0 GB/s	4.0 GB/s	8.0 GB/s
3.0	2010	128b/130b	8.0 GT/s	984.6 MB/s	1.97 GB/s	3.94 GB/s	7.88 GB/s	15.8 GB/s
4.0	2017	128b/130b	16.0 GT/s	1969 MB/s	3.94 GB/s	7.88 GB/s	15.75 GB/s	31.5 GB/s
5.0 ^{[32][33]}	expected in Q2 2019 ^[34]	128b/130b	32.0 GT/s ^[ii]	3938 MB/s	7.88 GB/s	15.75 GB/s	31.51 GB/s	63.0 GB/s

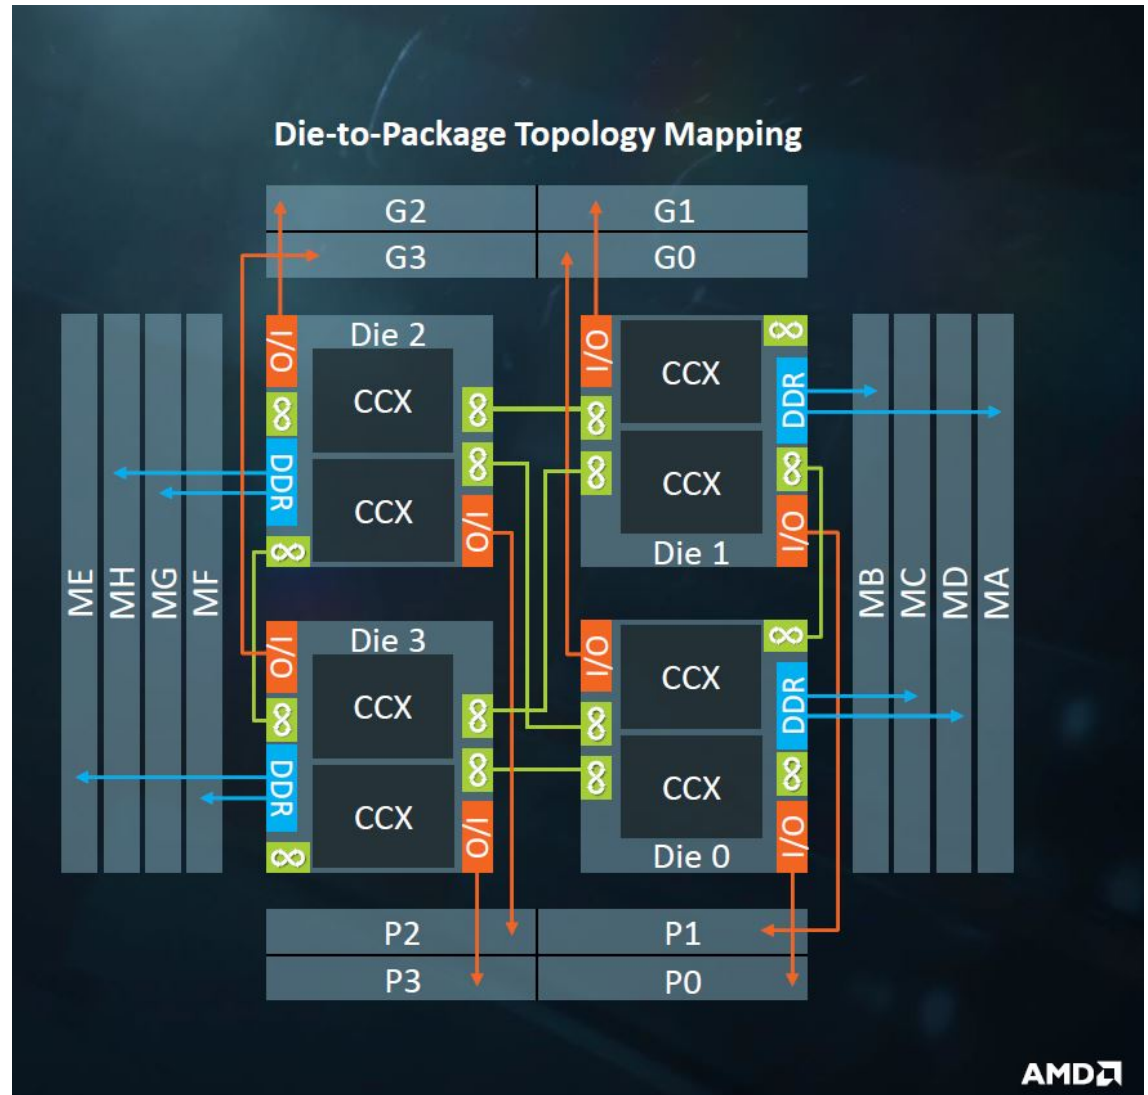
The PowerEdge R640(1U) system supports PCI express (PCIe) generation 3 expansion cards (4)

PCI buses on ORFEO epyc nodes


CCX is a core complex of up to 4 cores that share L3 cache. M* are the memory channels, two channels handled by each die. P* and G* are IO lanes. ∞ is the Infinity Fabric.

On a single-socket system, each die provides up to 32 PCI-E lanes using the P* and the G* IO lanes shown in Figure.

In total 128 IO lines



2 socket layout:



in a two-socket (2S) configuration, half the IO lanes of each die are used to connect to one of the dies on the other socket by using the G* IO lanes configured as Infinity Fabric. This leaves the socket with the P* IO lanes for a total of 64 PCI-E lanes and, thus, still 128 PCI-E lanes for the platform.

PCI buses on ORFEO epyc nodes

- Let us use hwloc and lstopo

- 8 NUMA regions




MPI libraries available on ORFEO

- Open-MPI
 - Open-source
 - Portable and efficient
- Intel MPI
 - closed source
 - Fits perfectly the intel architecture



Open MPI

- ❑ Evolution of several prior MPI's
- ❑ Open source project and community
- ❑ Production quality
- ❑ Vendor-friendly
- ❑ Research- and academic-friendly
- ❑ MPI-3.1 compliant

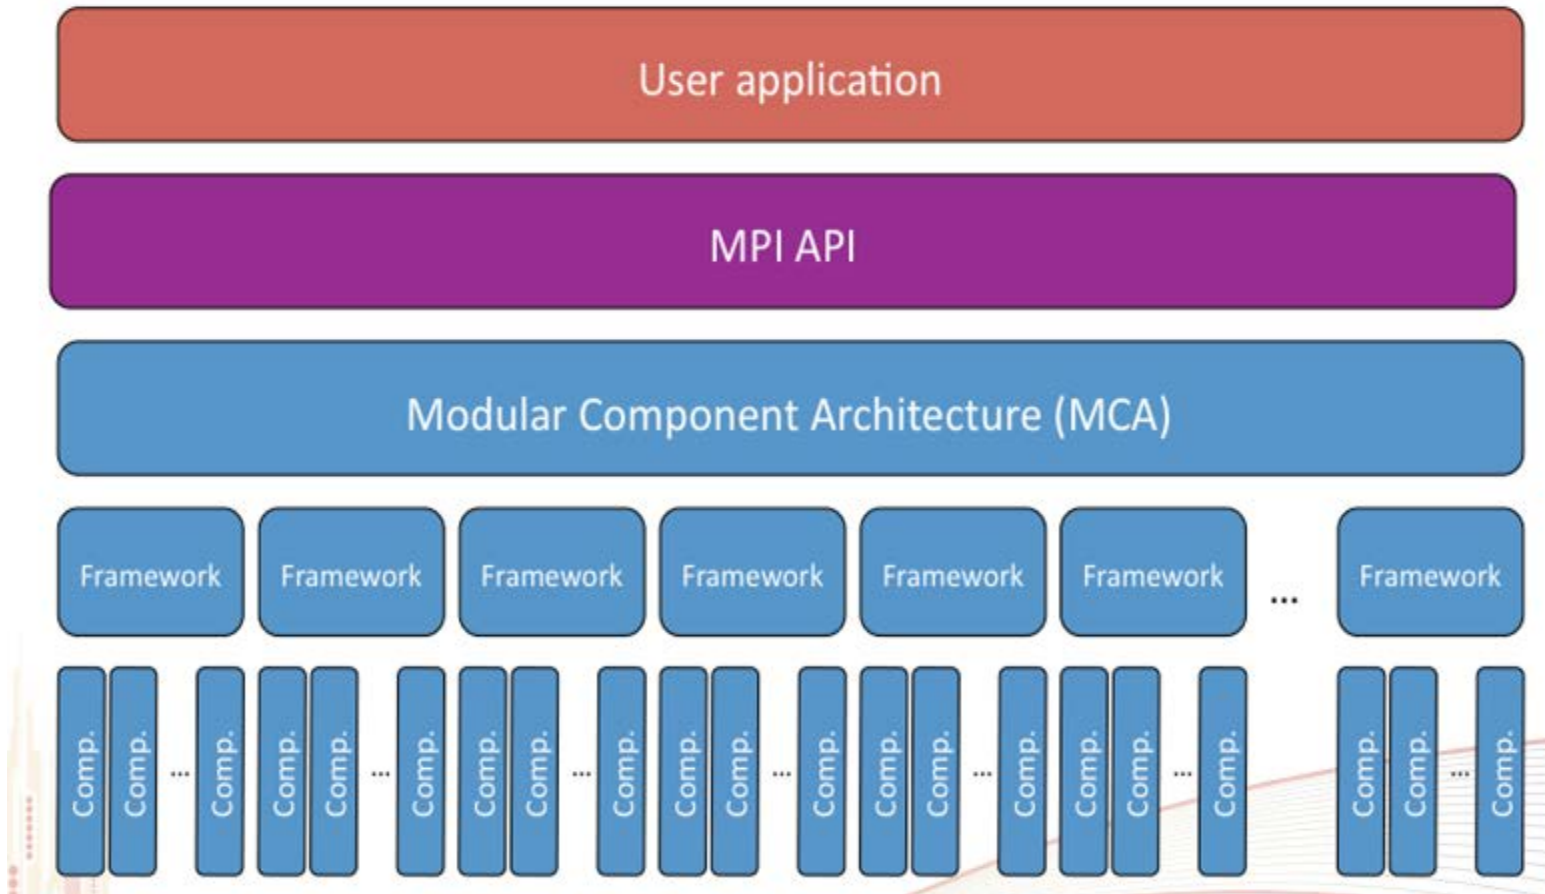


<https://www.open-mpi.org/>

openMPI: Plugins for HPC

- Uses Modular Component Architecture (MCA)
- Run-time plugins for combinatorial functionality
 - Underlying point-to-point network support
 - Different MPI collective algorithms
 - back-end run-time environment / scheduler support
- Extensive run-time tuning capabilities
 - Allow user or system administrator to tweak performance for a given platform

Plugin high level view



Lots and lots of plugin type

- Back-end network
- Resource manager support
- Operating system support
- All can be loaded (or not) at runtime
- Choice of network is a runtime decision

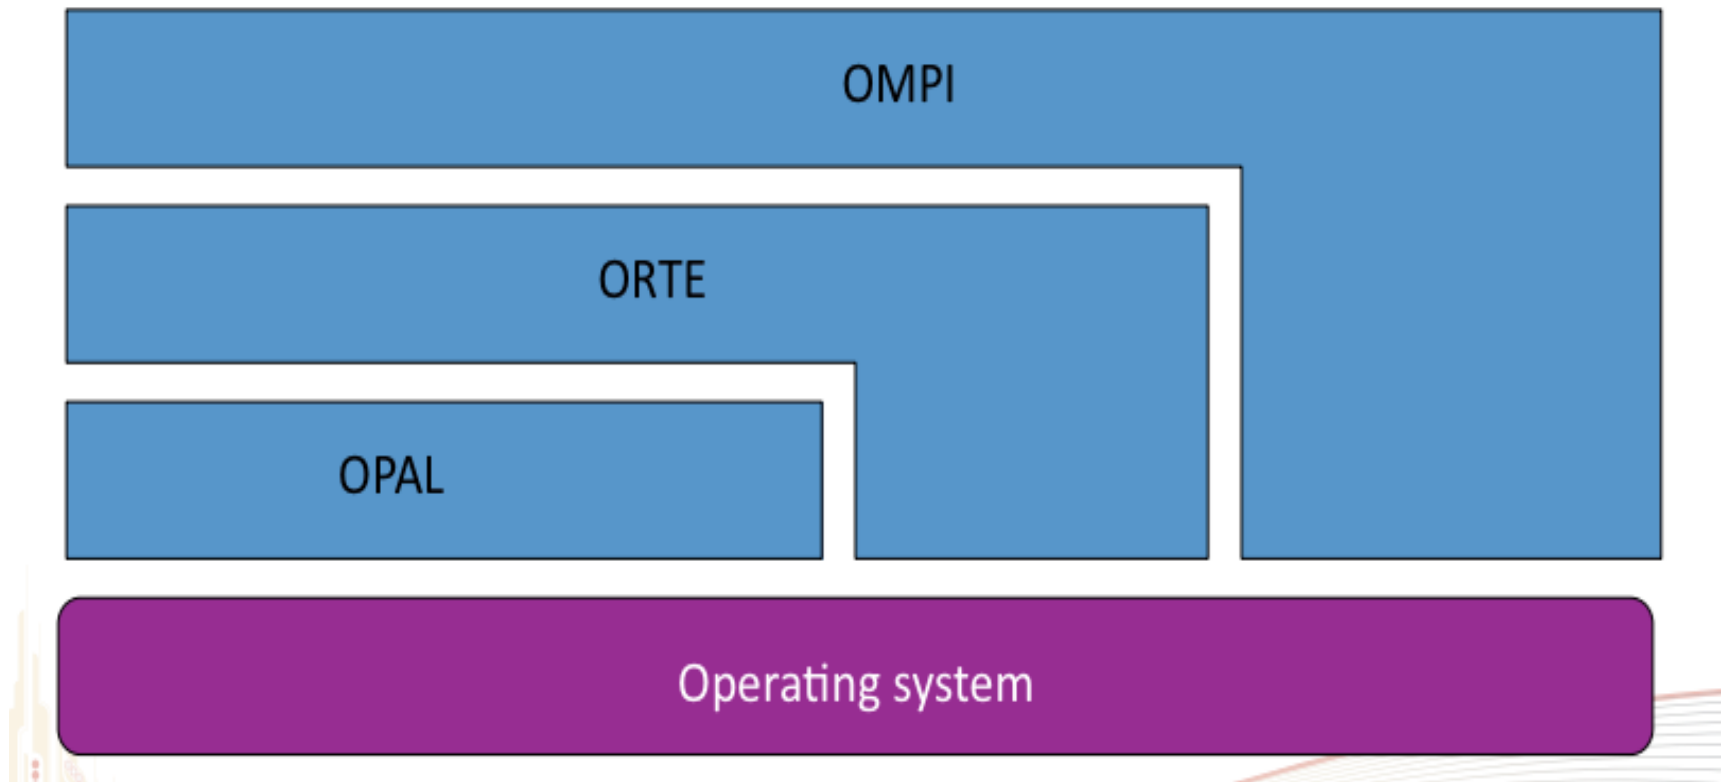
MPI frameworks (version 4.x)

- bml: BTL multipliexing layer
- btl: Byte transport layer
- coll: MPI collectives
- fbt: MPI file byte transfer layer
- fcoll: MPI file collectives
- fs: MPI file management
- hook: Generic hooks
- io: MPI IO
- mtl: Matching transport layer
- op: MPI reduction operations
- osc: MPI one sided communications
- pml: MPI point-to-point communications
- sharedfp: MPI shared file pointer operations
- topo: MPI topologies
- vprotocol: Virtual protocol API interposition

OpenMPI software stack

- Three main section:
- OpenMPI layer (OMPI)
- OpenMPI Run-Time environment (ORTE)
- Open Portability Access Layer (OPAL)
- OMPI → ORTE → OPAL

Graphical view



Ompi_info...

```
[cozzini@login02 ~]$ srun -n1 ompi_info | grep btl
srun: Warning: can't run 1 processes on 2 nodes, setting nnodes to 1
MCA btl: self (MCA v2.1.0, API v3.1.0, Component v4.1.4)
MCA btl: ofi (MCA v2.1.0, API v3.1.0, Component v4.1.4)
MCA btl: openib (MCA v2.1.0, API v3.1.0, Component v4.1.4)
MCA btl: tcp (MCA v2.1.0, API v3.1.0, Component v4.1.4)
MCA btl: usnic (MCA v2.1.0, API v3.1.0, Component v4.1.4)
MCA btl: vader (MCA v2.1.0, API v3.1.0, Component v4.1.4)
MCA fbtl: posix (MCA v2.1.0, API v2.0.0, Component v4.1.4)
[cozzini@login02 ~]$ srun -n1 ompi_info | grep pml
srun: Warning: can't run 1 processes on 2 nodes, setting nnodes to 1
MCA pml: v (MCA v2.1.0, API v2.0.0, Component v4.1.4)
MCA pml: cm (MCA v2.1.0, API v2.0.0, Component v4.1.4)
MCA pml: monitoring (MCA v2.1.0, API v2.0.0, Component v4.1.4)
MCA pml: ob1 (MCA v2.1.0, API v2.0.0, Component v4.1.4)
MCA pml: ucx (MCA v2.1.0, API v2.0.0, Component v4.1.4)
```

Point to Point component Frameworks

- Byte Transfer Layer (BTL)
 - Abstracts lowest native network interfaces
- Point-to-Point Messaging Layer (PML)
 - Implements MPI semantics, message fragmentation, and striping across BTLs

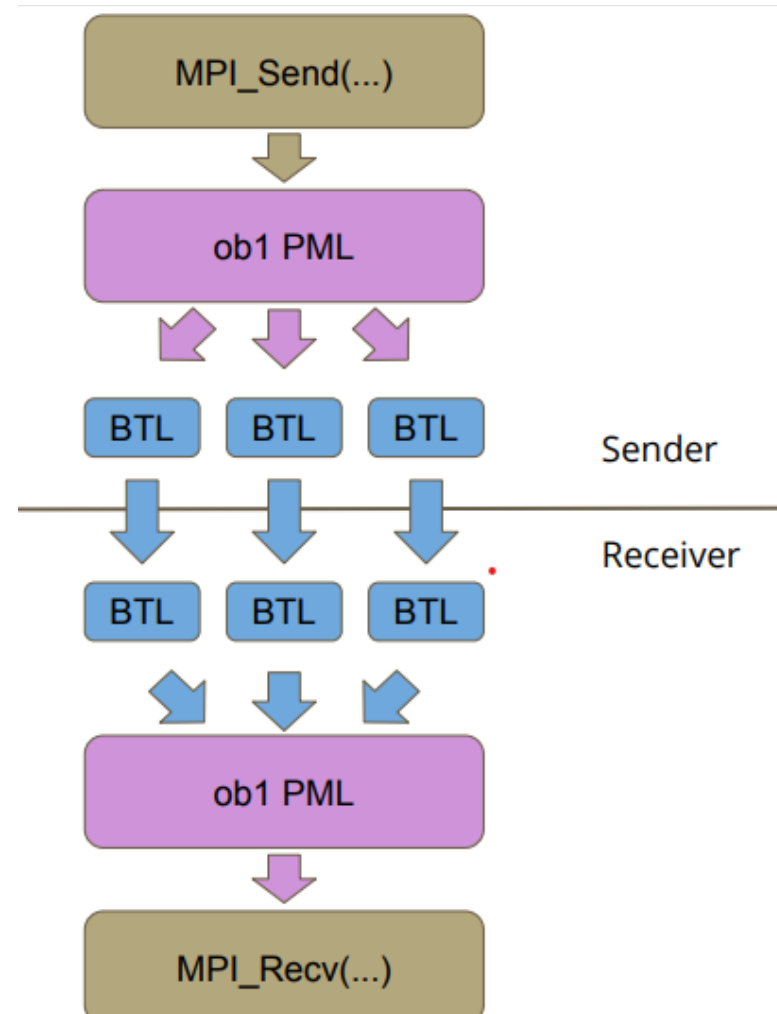
Shall we need to know all this ?

PML details..

- There are several PMLs to choose from:
 - ob1: Multi-device, multi-rail engine
 - Uses BTL components (byte transfer layer)
 - cm: Engine for matching network layers
 - Uses MTL components (matching transport layer)
 - ucx: Uses the UCX communication library (Unified Communications X)

ob1: Multi-Device, Multi-Rail Engine

- ob1 will:
 - 1. Pick BTL instance(s) that can reach a given peer
 - 2. Split large messages across relevant BTL instances
 - 3. Re-assemble messages at the receiver
- ob1 was Open MPI's original point-to-point transport engine and still works well in many environments.




Available BTLs..

- ofi: Libfabric (OpenFabrics Interfaces)
- portals4: Portals-based networks (uncommon)
- self: Process-loopback communications
- ~~sm~~ vader: Shared memory
- smcuda: CUDA-aware shared memory
- tcp: TCP
- uct: UCX
- ugni: Cray uGNI (userspace Generic Network Interface)
- usnic: Cisco usNIC (userspace NIC)

ucx: Thin Interface to the UCX Library

- UCX is, itself, a multi-device, multi-rail transport library. It has its own engine, and therefore did not need another engine in Open MPI.
- Hence, the UCX community decided to write their own (very thin) PML and not use an existing Open MPI engine.
- NOTE: The diagram only shows the MPI code blocks (not the UCX library itself)



By default, which network gets used at run time?

UCX for Infiniband


CML + PSM2 MTL For
OmniPath

OB1 PML + BTLs For
all others

self

vader

tcp



```
mpirun -np 2 ./IMB-MPI1 PingPong
```


What should I do to use a different network stack ?

- Force the use of OB1 and BTLs:

```
mpirun --mca pml ob1 --mca btl [comma-delimited list]
```

- Force the use of CM and MTLs:

```
mpirun --mca pml ob1 --mca btl [comma-delimited list]
```

- Force the use of the UCX PML:

```
mpirun --mca pml ucx
```

Using OpenMPI library with ob1

- Tests to perform:

```
mpirun -np 2 --mca pml ob1 --report-bindings --map-by  
node --mca btl tcp,self ./IMB-MPI1 PingPong
```

```
mpirun -np 2 --mca pml ob1 --report-bindings --map-by  
socket --mca btl tcp,self ./IMB-MPI1 PingPong
```

```
mpirun -np 2 --mca pml ob1 --report-bindings --map-by  
core --mca btl tcp,self ./IMB-MPI1 PingPong
```

See tutorial in MPI directory

MPI Intel library

- Intel® MPI Library is a multi-fabric message-passing library that implements the open-source MPICH specification.
- Highly tuned on HPC clusters based on Intel® processors.
 - Achieve the best latency, bandwidth, and scalability through automatic tuning for the latest Intel® platforms.
- Fully integrated in the Intel Cluster edition with compilers, math libraries and performance tuner/analyzer..

Using Intel MPI

- Load module
- Check wrapper:

```
[cozzini@ctlpg-gnode001 src_c-intel]$ mpicc -v
mpigcc for the Intel(R) MPI Library 2019 Update 9 for Linux*
Copyright 2003-2020, Intel Corporation.
Using built-in specs.
COLLECT_GCC=gcc
...
Thread model: posix
gcc version 4.8.5 20150623 (Red Hat 4.8.5-39) (GCC)
[cozzini@ctlpg-gnode001 src_c-intel]$ mpiicc -v
mpiicc for the Intel(R) MPI Library 2019 Update 9 for Linux*
Copyright 2003-2020, Intel Corporation.
icc version 19.1.3.304 (gcc version 4.8.5 compatibility)
```

Using Intel MPI

- Run benchmark:

```
cozzini@ctlpg-gnode00 ]$ which mpirun
/opt/area/shared/programs/x86_64/intel/parallel_studio_xe_2020_update4_cluster_edition/compilers_and_libraries_2020/linux/mpi/intel64/bin/mpirun

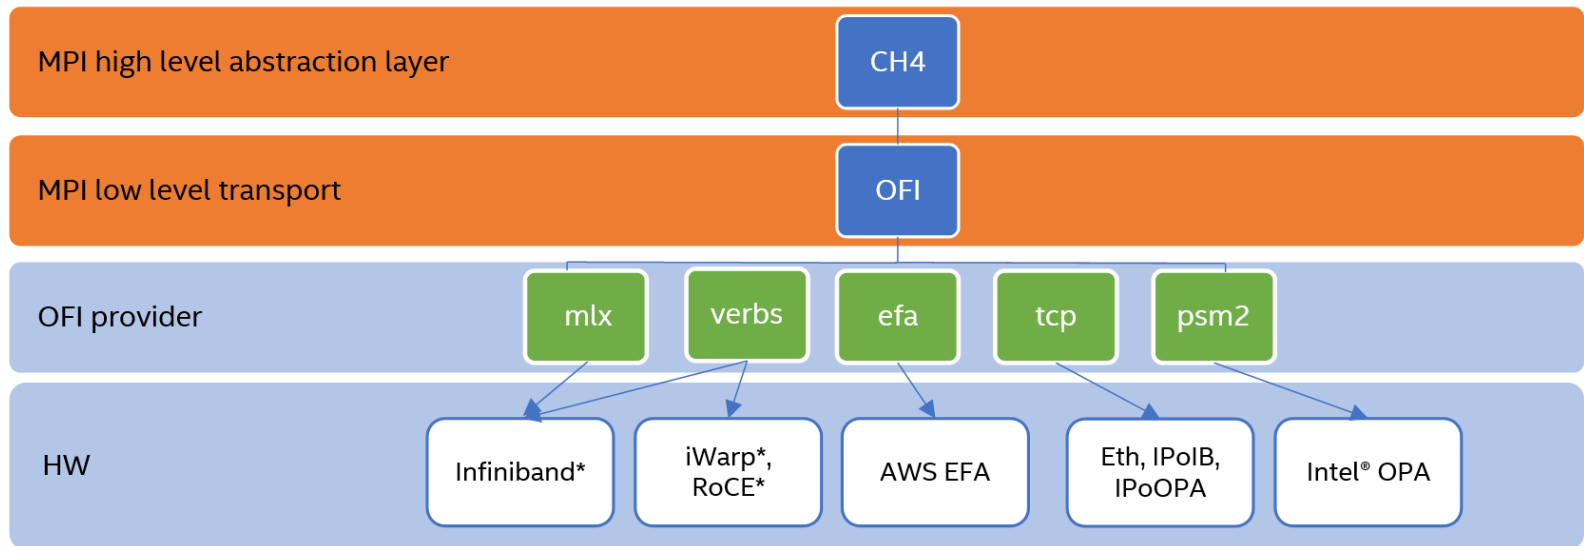
[cozzini@ctlpg-gnode001]$ mpirun -np 2 ./IMB-MPI1 PingPong -msglog 4
...
#-----
# Benchmarking PingPong
# #processes = 2
#-----
#
```

#bytes	#repetitions	t[usec]	Mbytes/sec
0	1000	0.47	0.00
1	1000	0.47	2.12
2	1000	0.47	4.26
4	1000	0.47	8.52
8	1000	0.47	17.06
16	1000	0.47	34.10

```
# All processes entering MPI_Finalize
```

Tuning Intel MPI

- Intel MPI Library software stack:



- The Intel® MPI Library will attempt to select the fastest available fabric by default,

Tuning Intel MPI

- Pinning MPI process on specific processors:

I_MPI_PIN_PROCESSOR_LIST generates a custom process to processor map with one of the three alternative syntax available:

<proclist>

<procset>[:map=<map>]

<procset>[:[grain=<grain>][,shift=<shift>][,preoffset=<preoffset>][,postoffset=<postoffset>]]

Some examples:

- Run on the two contiguous processors:

```
[cozzini@ctlpg-gnode001 src_c-intel]$ mpirun -np 2 -ppn=2 -env I_MPI_DEBUG 5 -genv I_MPI_PIN_PROCESSOR_LIST 0,1 ./IMB-MPI1 PingPong -msglog 4
[0] MPI startup(): Intel(R) MPI Library, Version 2019 Update 9 Build 20200923 (id: abd58e492)
[0] MPI startup(): Copyright (C) 2003-2020 Intel Corporation. All rights reserved.
[0] MPI startup(): library kind: release
[0] MPI startup(): libfabric version: 1.10.1-impi
[0] MPI startup(): libfabric provider: mlx
[0] MPI startup(): Rank      Pid      Node name      Pin cpu
[0] MPI startup(): 0        32018   ctlpg-gnode001  0
[0] MPI startup(): 1        32019   ctlpg-gnode001  1
[0] MPI startup():
I_MPI_ROOT=/opt/area/shared/programs/x86_64/intel/parallel_studio_xe_2020_update4_cluster_edition//compilers_and_libraries_2020/linux/mpi
[0] MPI startup(): I_MPI_MPIRUN=mpirun
[0] MPI startup(): I_MPI_HYDRA_RMK=pbs
[0] MPI startup(): I_MPI_HYDRA_TOPOLIB=hwloc
[0] MPI startup(): I_MPI_PIN_PROCESSOR_LIST=0,1
[0] MPI startup(): I_MPI_INTERNAL_MEM_POLICY=default
[0] MPI startup(): I_MPI_DEBUG=5

#-----
# Benchmarking PingPong
# #processes = 2
#-----


| #bytes | #repetitions | t[usec] | Mbytes/sec |
|--------|--------------|---------|------------|
| 0      | 1000         | 0.47    | 0.00       |
| 1      | 1000         | 0.47    | 2.11       |
| 2      | 1000         | 0.47    | 4.22       |
| 4      | 1000         | 0.47    | 8.44       |
| 8      | 1000         | 0.47    | 16.87      |
| 16     | 1000         | 0.47    | 33.90      |


```


Some examples:

- Run on the on the same socket:

```
[cozzini@ctlpg-gnode001 src_c-intel]$ mpirun -np 2 -env I_MPI_DEBUG 5 -genv I_MPI_PIN_PROCESSOR_LIST 0,2
./IMB-MPI1 PingPong
[0] MPI startup(): Intel(R) MPI Library, Version 2019 Update 9 Build 20200923 (id: abd58e492)
[0] MPI startup(): Copyright (C) 2003-2020 Intel Corporation. All rights reserved.
[0] MPI startup(): library kind: release
[0] MPI startup(): libfabric version: 1.10.1-impi
[0] MPI startup(): libfabric provider: mlx
[0] MPI startup(): Rank      Pid      Node name      Pin cpu
[0] MPI startup(): 0        32479   ctlpg-gnode001 0
[0] MPI startup(): 1        32480   ctlpg-gnode001 2
[0] MPI startup():
I_MPI_ROOT=/opt/area/shared/programs/x86_64/intel/parallel_studio_xe_2020_update4_cluster_edition//compilers
_and_libraries_2020/linux/mpi
[0] MPI startup(): I_MPI_MPIRUN=mpirun
[0] MPI startup(): I_MPI_HYDRA_RMK=pbs
[0] MPI startup(): I_MPI_HYDRA_TOPOLIB=hwloc
[0] MPI startup(): I_MPI_PIN_PROCESSOR_LIST=0,2
[0] MPI startup(): I_MPI_INTERNAL_MEM_POLICY=default
[0] MPI startup(): I_MPI_DEBUG=5
....
# Benchmarking PingPong
# #processes = 2
#-----
#bytes #repetitions      t[usec]      Mbytes/sec
      0           1000         0.28         0.00
      1           1000         0.29         3.47
      2           1000         0.29         6.91
      4           1000         0.29        13.84
      8           1000         0.29        27.84
```

Comparing performance

- Left to the readers... 😊

- A starting point for intel:

[Tuning the Intel® MPI Library: Basic Techniques](#)

Final considerations 1

- Why latency is so important ?
- According to Amdahl's law:
 - a high-performance parallel system tends to be bottlenecked by its slowest sequential process
- in all but the most embarrassingly parallel supercomputer workloads, the slowest sequential process is often the latency of message transmission across the network

Final considerations 2

- In general the compute/communication ratio in a parallel program remains fairly constant.
- So as the computational power increases the network speed must also be increased.
- We are living in a multi-core world: MPI processes **sharing the same network device !**
- Contention for the interconnect device can have a significant impact on performance.