

Foundations of High Performance Computing

Using High Performance Libraries



Agenda

- Intro:
 - What are we using HPC for ?
 - Where should we start optimizing ?
- High Performance Libraries
- Linear Algebra libraries
- Using HP libraries: some examples.

The seven dwarfs of HPC



- Phil Colella (LBL) identified **7 kernels** of which most simulation and data analysis program are composed:
- Dense Linear Algebra
 - Ex: solve $Ax=B$ or $Ax=\lambda x$ where A is a dense matrix
- Sparse Linear Algebra
 - solve $Ax=B$ or $Ax=\lambda x$ where A is a sparse matrix (mostly zero)
- Operation on structured Grids:
 - $ANEW_j() = 4 * (A(i,j) - A(i-1,j) - A(i+1,j) - A(i,j-1) - A(i,j+1))$
- Operation on unstructured Grids:
 - similar but list of neighbours varies from entry to entry
- Spectral Methods
 - Fast Fourier Transform (FFT)
- Particle Methods
 - Compute electrostatic forces on n -particles
- Monte Carlo
 - many independent simulation using different inputs

Is this the real picture in 2022 ?

We are missing ALL data ML/DL working load..


But...

A lot of them relies on highly optimized libraries.

Where should you start
optimizing your application ?

Optimization techniques

- There are basically three different categories:
 - Improve memory performance (the most important)
 - Improve CPU performance
 - Use already highly optimized libraries/subroutines



The easiest and more efficient way..

What are High Performance Libraries ?

- Routines for common (math) functions written in a specific way to take advantage of all capabilities of the CPU.
- Each CPU type normally has its own version of the library specifically written or compiled to maximally exploit that architecture

Why using High Performance Libraries ?

- Compilers can optimize code only to a certain point. Effective programming needs deep knowledge of the platform
- Performance libraries are designed to use the CPU in the most efficient way, which is not necessarily the most straightforward way.
- It is normally best to use the libraries supplied by or recommended by the CPU vendor
- On modern hardware they are hugely important, as they most efficiently exploit caches, special instructions and parallelism
- **Parallelism (at least on single node) comes for free..**

Any other reason apart from performance ?


- Usage of libraries makes coding easier.
Complicated math operations can be used from existing routines
- Increase portability of code as standard (and well optimized) libraries exist for ALL computing platforms.
- Lego approach: build your own code using already available bricks..

What is available ?

- Linear Algebra:
 - BLAS/LAPACK/SCALAPACK
- FFT:
 - FFTW
- ODE/PDE
 - PETSC
- Machine Learning:
 - Tensorflow / Caffe etc..

Should I write my own algorithm for L. A. ?

- 99.99% of time: NO !
- Tons of libraries out there
- Well tested
- Extremely efficient in 99.99% of the case
- With some “de facto” standard implemented



PORTABILITY IS COMING (almost) FOR FREE

Why dense Linear Algebra ?

- Many large matrices are sparse, but ...
- Dense algorithms easier to understand
- Some applications yields large dense matrices
- Large sparse matrix algorithms often yield smaller (but still large) dense problems
- LINPACK Benchmark (www.top500.org)

“ How fast is your computer?”

=

“How fast can you solve dense $Ax=b$?”

BLAS

Basic Linear Algebra Subprograms

BLAS summary

Basic Linear Algebra Subroutines

| Name | Description | Examples |
|--------------|--------------------------|---------------------------------|
| Level-1 BLAS | Vector Operations | $C = \sum X_i Y_i$ |
| Level-2 BLAS | Matrix-Vector Operations | $B_i = \sum_k A_{ik} X_k$ |
| Level-3 BLAS | Matrix-Matrix Operations | $C_{ij} = \sum_k A_{ik} B_{kj}$ |

BLAS list

Level 1 BLAS

| | dim | scalar | vector | vector | scalars | 5-element array |
|-----------------------|--------|--------|--------------------|--------|---------------|-----------------|
| SUBROUTINE xROTG (| | | | | A, B, C, S) | |
| SUBROUTINE xROTMG(| | | | | D1, D2, A, B, | PARAM) |
| SUBROUTINE xROT (N, | | | I, INCX, Y, INCY, | | C, S) | PARAM) |
| SUBROUTINE xROTM (N, | | | I, INCX, Y, INCY, | | | |
| SUBROUTINE xSWAP (N, | | | I, INCX, Y, INCY) | | | |
| SUBROUTINE xSCAL (N, | ALPHA, | | I, INCX) | | | |
| SUBROUTINE xCOPY (N, | | | I, INCX, Y, INCY) | | | |
| SUBROUTINE xAXPY (N, | ALPHA, | | I, INCX, Y, INCY) | | | |
| FUNCTION xDOT (N, | | | I, INCX, Y, INCY) | | | |
| FUNCTION xDOTU (N, | | | I, INCX, Y, INCY) | | | |
| FUNCTION xDOTC (N, | | | I, INCX, Y, INCY) | | | |
| FUNCTION xxDOT (N, | | | I, INCX, Y, INCY) | | | |
| FUNCTION xRN2 (N, | | | I, INCX) | | | |
| FUNCTION xASUM (N, | | | I, INCX) | | | |
| FUNCTION IxAMAX(N, | | | I, INCX) | | | |

Generate plane rotation
 Generate modified plane rotation
 Apply plane rotation
 Apply modified plane rotation
 $x \leftrightarrow y$
 $x \leftarrow \alpha x$
 $y \leftarrow x$
 $y \leftarrow \alpha x + y$
 $dot \leftarrow x^T y$
 $dot \leftarrow x^T y$
 $dot \leftarrow x^H y$
 $dot \leftarrow \alpha + x^T y$
 $rnrm2 \leftarrow ||x||_2$
 $asum \leftarrow ||re(x)||_1 + ||im(x)||_1$
 $amax \leftarrow 1^{*}k \ni |re(x_k)| + |im(x_k)|$
 $= \max\{|re(x_i)| + |im(x_i)|\}$

prefixes
 S, D
 S, D
 S, D
 S, D
 S, D, C, Z
 S, D, C, Z, CS, ZD
 S, D, C, Z
 S, D, C, Z
 S, D, C, Z
 S, D, DS
 C, Z
 C, Z
 SDS
 S, D, SC, DZ
 S, D, SC, DZ
 S, D, C, Z

Level 2 BLAS

| options | dim | b-width | scalar | matrix | vector | scalar | vector |
|----------------------------|---------------|---------|----------|-----------|-----------|--------|-----------|
| xGEMV (TRANS, | M, N, | | ALPHA, | A, LDA, | X, INCX, | BETA, | Y, INCY) |
| xGEMV (TRANS, | M, N, KL, KU, | | ALPHA, | A, LDA, | X, INCX, | BETA, | Y, INCY) |
| xHEMV (UPLO, | N, | | ALPHA, | A, LDA, | X, INCX, | BETA, | Y, INCY) |
| xHEMV (UPLO, | N, K, | | ALPHA, | A, LDA, | X, INCX, | BETA, | Y, INCY) |
| xHPMV (UPLO, | N, | | ALPHA, | AP, | X, INCX, | BETA, | Y, INCY) |
| xSYMV (UPLO, | N, | | ALPHA, | A, LDA, | X, INCX, | BETA, | Y, INCY) |
| xHEMV (UPLO, | N, K, | | ALPHA, | A, LDA, | X, INCX, | BETA, | Y, INCY) |
| xSPMV (UPLO, | N, | | ALPHA, | AP, | X, INCX, | BETA, | Y, INCY) |
| xTRMV (UPLO, TRANS, DIAG, | N, | | A, | LDA, | X, INCX) | | |
| xTRMV (UPLO, TRANS, DIAG, | N, K, | | A, | LDA, | X, INCX) | | |
| xTRMV (UPLO, TRANS, DIAG, | N, | | AP, | X, INCX) | | | |
| xTRSV (UPLO, TRANS, DIAG, | N, | | A, | LDA, | X, INCX) | | |
| xTRSV (UPLO, TRANS, DIAG, | N, K, | | A, | LDA, | X, INCX) | | |
| xTPSV (UPLO, TRANS, DIAG, | N, | | AP, | X, INCX) | | | |
| options | dim | scalar | vector | vector | matrix | | |
| xGER (| M, N, | ALPHA, | X, INCX, | Y, INCY, | A, LDA) | | |
| xGERU (| M, N, | ALPHA, | X, INCX, | Y, INCY, | A, LDA) | | |
| xGERC (| M, N, | ALPHA, | X, INCX, | Y, INCY, | A, LDA) | | |
| xHER (UPLO, | N, | ALPHA, | X, INCX, | | A, LDA) | | |
| xHPR (UPLO, | N, | ALPHA, | X, INCX, | AP) | | | |
| xHER2 (UPLO, | N, | ALPHA, | X, INCX, | Y, INCY, | A, LDA) | | |
| xHPR2 (UPLO, | N, | ALPHA, | X, INCX, | Y, INCY, | AP) | | |
| xSYR (UPLO, | N, | ALPHA, | X, INCX, | | A, LDA) | | |
| xSPR (UPLO, | N, | ALPHA, | X, INCX, | | AP) | | |
| xSYR2 (UPLO, | N, | ALPHA, | X, INCX, | Y, INCY, | A, LDA) | | |
| xSPR2 (UPLO, | N, | ALPHA, | X, INCX, | Y, INCY, | AP) | | |

$y \leftarrow \alpha Ax + \beta y, y \leftarrow \alpha A^T x + \beta y, y \leftarrow \alpha A^H x + \beta y, A - m \times n$
 $y \leftarrow \alpha Ax + \beta y, y \leftarrow \alpha A^T x + \beta y, y \leftarrow \alpha A^H x + \beta y, A - m \times n$
 $y \leftarrow \alpha Ax + \beta y$
 $y \leftarrow \alpha Ax + \beta y$
 $y \leftarrow \alpha Ax + \beta y$
 $y \leftarrow \alpha Ax + \beta y$
 $y \leftarrow \alpha Ax + \beta y$
 $y \leftarrow \alpha Ax + \beta y$
 $x \leftarrow Ax, x \leftarrow A^T x, x \leftarrow A^H x$
 $x \leftarrow Ax, x \leftarrow A^T x, x \leftarrow A^H x$
 $x \leftarrow Ax, x \leftarrow A^T x, x \leftarrow A^H x$
 $x \leftarrow A^{-1} x, x \leftarrow A^{-T} x, x \leftarrow A^{-H} x$
 $x \leftarrow A^{-1} x, x \leftarrow A^{-T} x, x \leftarrow A^{-H} x$
 $x \leftarrow A^{-1} x, x \leftarrow A^{-T} x, x \leftarrow A^{-H} x$

S, D, C, Z
 S, D, C, Z
 C, Z
 C, Z
 C, Z
 S, D
 S, D
 S, D, C, Z
 S, D, C, Z
 S, D, C, Z
 S, D, C, Z
 S, D, C, Z
 S, D, C, Z
 S, D, C, Z

$A \leftarrow \alpha xy^T + A, A - m \times n$
 $A \leftarrow \alpha xy^T + A, A - m \times n$
 $A \leftarrow \alpha xy^H + A, A - m \times n$
 $A \leftarrow \alpha xx^H + A$
 $A \leftarrow \alpha xx^H + A$
 $A \leftarrow \alpha xy^H + y(\alpha x)^H + A$
 $A \leftarrow \alpha xy^H + y(\alpha x)^H + A$
 $A \leftarrow \alpha xx^T + A$
 $A \leftarrow \alpha xx^T + A$
 $A \leftarrow \alpha xy^T + \alpha yx^T + A$
 $A \leftarrow \alpha xy^T + \alpha yx^T + A$

S, D
 C, Z
 C, Z
 C, Z
 C, Z
 C, Z
 C, Z
 S, D
 S, D
 S, D
 S, D

Level 3 BLAS

| options | dim | scalar | matrix | matrix | scalar | matrix |
|-----------------------------|-------------|--------|---------|----------|--------|----------|
| xGEMM (TRANSA, TRANSB, | M, N, K, | ALPHA, | A, LDA, | B, LDB, | BETA, | C, LDC) |
| xSYMM (SIDE, UPLO, | M, N, | ALPHA, | A, LDA, | B, LDB, | BETA, | C, LDC) |
| xHEMM (SIDE, UPLO, | M, N, | ALPHA, | A, LDA, | B, LDB, | BETA, | C, LDC) |
| xSYRK (UPLO, TRANS, | N, K, | ALPHA, | A, LDA, | | BETA, | C, LDC) |
| xHERK (UPLO, TRANS, | N, K, | ALPHA, | A, LDA, | | BETA, | C, LDC) |
| xSYR2X (UPLO, TRANS, | N, K, | ALPHA, | A, LDA, | B, LDB, | BETA, | C, LDC) |
| xHER2X (UPLO, TRANS, | N, K, | ALPHA, | A, LDA, | B, LDB, | BETA, | C, LDC) |
| xTRMM (SIDE, UPLO, TRANSA, | DIAG, M, N, | ALPHA, | A, LDA, | B, LDB) | | |
| xTRSM (SIDE, UPLO, TRANSA, | DIAG, M, N, | ALPHA, | A, LDA, | B, LDB) | | |

$C \leftarrow \alpha op(A) op(B) + \beta C, op(X) = X, X^T, X^H, C - m \times n$
 $C \leftarrow \alpha AB + \beta C, C \leftarrow \alpha BA + \beta C, C - m \times n, A = A^T$
 $C \leftarrow \alpha AB + \beta C, C \leftarrow \alpha BA + \beta C, C - m \times n, A = A^H$
 $C \leftarrow \alpha AA^T + \beta C, C \leftarrow \alpha A^T A + \beta C, C - n \times n$
 $C \leftarrow \alpha AA^H + \beta C, C \leftarrow \alpha A^H A + \beta C, C - n \times n$
 $C \leftarrow \alpha AB^T + \delta BA^T + \beta C, C \leftarrow \alpha A^T B + \delta B^T A + \beta C, C - n \times n$
 $C \leftarrow \alpha AB^H + \delta BA^H + \beta C, C \leftarrow \alpha A^H B + \delta B^H A + \beta C, C - n \times n$
 $B \leftarrow \alpha op(A) B, B \leftarrow \alpha B op(A), op(A) = A, A^T, A^H, B - m \times n$
 $B \leftarrow \alpha op(A^{-1}) B, B \leftarrow \alpha B op(A^{-1}), op(A) = A, A^T, A^H, B - m \times n$

S, D, C, Z
 S, D, C, Z
 C, Z
 S, D, C, Z
 C, Z
 S, D, C, Z
 C, Z
 S, D, C, Z
 S, D, C, Z
 S, D, C, Z

Where can I get BLAS?

www.netlib.org/blas

- Source: 142 routines, 31K LOC,
- Testing: 28K LOC
- Reference (unoptimized) implementation only !
- http://www.netlib.org/blas/#_reference_blas_version_3_11_0
- Ex: 3 nested loops for GEMM

Why BLAS are important ?

- Because the BLAS are **efficient**, **portable**, **parallel**, and **widely available**, they are commonly used in the development of high quality linear algebra software.
- Performance of lot of applications depends a lot on the performance of the underlying BLAS
- Lot of applications include ML/DL stuff as well....
 - <https://petewarden.com/2015/04/20/why-gemm-is-at-the-heart-of-deep-learning/>

Standardization: BLAS example

- Each BLAS Subroutines have a standardized layout
- BLAS is documented in the source code
- Man pages exist
- Vendor supplied docs
- Different BLAS implementations have the same calling sequence

```
SUBROUTINE DGEMM ( TRANSA, TRANSB, N, N, K, ALPHA, A, LDA, B, LDB,
#      BETA, C, LDC )
*
* .. SCALAR ARGUMENTS ..
* CHARACTER*1      TRANSA, TRANSB
* INTEGER          M, N, K, LDA, LDB, LDC
* DOUBLE PRECISION ALPHA, BETA
* .. ARRAY ARGUMENTS ..
* DOUBLE PRECISION A( LDA, * ), B( LDB, * ), C( LDC, * )
*
* ..
*
* PURPOSE
* =====
*
* DGEMM PERFORMS ONE OF THE MATRIX-MATRIX OPERATIONS
*
*   C := ALPHA*OP( A )*OP( B ) + BETA*C,
*
* WHERE OP( X ) IS ONE OF
*
*   OP( X ) = X   OR   OP( X ) = X',
*
* ALPHA AND BETA ARE SCALARS, AND A, B AND C ARE MATRICES, WITH OP( A )
* AN M BY K MATRIX, OP( B ) A K BY N MATRIX AND C AN M BY N MATRIX.
*
* PARAMETERS
* =====
*
* TRANSA - CHARACTER*1,
* ON ENTRY, TRANSA SPECIFIES THE FORM OF OP( A ) TO BE USED IN
* THE MATRIX MULTIPLICATION AS FOLLOWS:
*
*   TRANSA = 'N' OR 'n', OP( A ) = A,
*   TRANSA = 'T' OR 't', OP( A ) = A',
*   TRANSA = 'C' OR 'c', OP( A ) = A',
*
* UNCHANGED ON EXIT.
*
* TRANSB - CHARACTER*1,
* ON ENTRY, TRANSB SPECIFIES THE FORM OF OP( B ) TO BE USED IN
* THE MATRIX MULTIPLICATION AS FOLLOWS:
*
*   TRANSB = 'N' OR 'n', OP( B ) = B,
*   TRANSB = 'T' OR 't', OP( B ) = B',
```

Vendor/Optimized BLAS libraries

- ACML
 - The AMD Core Math Library, supporting the AMD processors
- ATLAS
 - Automatically Tuned Linear Algebra, an open source implementation of BLAS APIs for C and Fortran 77
- Intel MKL
 - The Intel Math Kernel Library supporting x86 32-bits and 64-bits. Includes optimizations for Intel Pentium, Core and Intel Xeon CPUs and Intel Xeon Phi; support for Linux, Windows and Mac OS X
- cuBLAS
 - Optimized BLAS for NVIDIA based GPU cards
- ESSL
 - IBM's Engineering and Scientific Subroutine Library, supporting the PowerPC architecture under AIX and Linux
- GotoBLAS
 - Kazushige Goto's BSD-licensed implementation of BLAS, tuned in particular for Intel, VIA Nanoprocessor, AMD Opteron
- OpenBLAS
 - Optimized BLAS based on Goto BLAS hosted at GitHub, supporting Intel platform and other
- And many others...

What about C/C++ program?

- BLAS routines are Fortran-style, when calling them from C language programs, follow the Fortran-style calling conventions:
 - Pass variables by address, not by value.
 - Store your data in Fortran style, that is, column-major rather than row-major order.
- Be aware that because the Fortran language is case-insensitive, the routine names can be both upper-case or lower-case, with or without the trailing underscore.
- For example, the following names are equivalent:
 - dgemm, DGEMM, dgemm_, and DGEMM_

Use CBLAS !

- C-style interface to the BLAS routines

www.netlib.org/blas/blast-forum/cblas.tgz

- You can call CBLAS routines using regular C- style calls.
- The header file specifies enumerated values and prototypes of all the functions.
- Details and examples here:

<https://software.intel.com/en-us/mkl-tutorial-c-multiplying-matrices-using-dgemm>

Q parameter: aka computational efficiency...

Table 2: Basic Linear Algebra Subroutines (BLAS)

| Operation | Definition | Floating point operations | Memory references | q |
|---------------------------|---|---------------------------|-------------------|-------|
| saxpy | $y_i = \alpha x_i + y_i, i = 1, \dots, n$ | $2n$ | $3n + 1$ | $2/3$ |
| Matrix-vector mult | $y_i = \sum_{j=1}^n A_{ij}x_j + y_i$ | $2n^2$ | $n^2 + 3n$ | 2 |
| Matrix-matrix mult | $C_{ij} = \sum_{k=1}^n A_{ik}B_{kj} + C_{ij}$ | $2n^3$ | $4n^2$ | $n/2$ |


The parameter q is the ratio of flops to memory references.
Generally:

1. Larger values of q maximize useful work to time spent moving data.
2. The higher the level of the BLAS, the larger q .

It follows...

- BLAS1 are **memory bounded** !(for each computation a memory transfer is required)
- BLAS2 are not so memory bounded (can have good performance on super-scalar architecture)
- BLAS3 can be very efficient on super-scalar computers because **not memory bounded**

BLAS performance on SandyBridge (1core)



Proposed exercise/tutorial

- Create the same graph for ORFEO cores using MKL and OpenBLAS
- STEPS:
 - Install OpenBLAS libraries in your directory
 - Write a small program to call the three routines
 - dgemm/dgemv/daxypi
 - Write a script to collect all sizes of interest
 - Make nice plots

Linking optimized libraries...

- OpenBLAS:

- It comes with cblas bundled in so no problem with C/C++
- Automatically includes lapack reference implementation
- Compilation is straightforward:

```
gcc -o test test.c -I  
/your_path/OpenBLAS/include/  
-L/your_path/OpenBLAS/lib -lopenblas
```

- MKL

- Generally complex and highly dependent on version and/or HW/SW implementation
- <https://software.intel.com/en-us/articles/intel-mkl-link-line-advisor>

Are these libraries multithreaded ?

- MKL
 - Both sequential and multithreaded version available

Intel® oneAPI Math Kernel Library (oneMKL) Link Line Advisor v6.13

Reset

| | |
|---------------------------------------|---|
| Select Intel® product: | Intel(R) Parallel Studio XE 2020 ▼ |
| Select OS: | Linux* ▼ |
| Select compiler: | GNU C/C++ ▼ |
| Select architecture: | Intel(R) 64 ▼ |
| Select dynamic or static linking: | Dynamic ▼ |
| Select interface layer: | C/Fortran API with 32-bit integer ▼ |
| Select threading layer: | OpenMP threading ▼ |
| Select OpenMP library: | <Select threading> OpenMP threading Sequential TBB threading |
| Enable OpenMP offload feature to GPU: | |

MKL: how to control number of threads?

- OpenMP threading ? → OMP_NUM_THREADS
- Other threading ? → MKL_NUM_THREADS
- Define yourself the number of threads:
 - Place `mkl_set_num_thread(N)` routine in your code.
- All MKL routines call takes precedence over any environment variables. and MKL environment Variables will take precedence over the OpenMP* environments.

More details here:

<https://software.intel.com/content/www/us/en/develop/articles/recommended-settings-for-calling-intel-mkl-routines-from-multi-threaded-applications.html>

OpenBLAS:

- By default : multithreaded version, maximum number of threads established by cores available on the machine when compilation is performed;
- An example (ORFEO BLAS):

```
OpenBLAS build complete. (BLAS CBLAS LAPACK LAPACKE)
```

```
OS          ... Linux
Architecture ... x86_64
BINARY      ... 64bit
C compiler  ... GCC (cmd & version : cc (GCC) 4.8.5 20150623 (Red Hat 4.8.5-39))
Fortran compiler ... GFORTRAN (cmd & version : GNU Fortran (GCC) 9.3.0)
Library Name ... libopenblas_haswellp-r0.3.13.a (Multi-threading; Max num-threads is 48)
```

```
To install the library, you can run "make PREFIX=/path/to/your/installation install".
```

OpenBLAS: caveat

- If your application is already multi-threaded, **it will conflict** with OpenBLAS multi-threading. Thus, you must set OpenBLAS to use single thread as following.
 - export OPENBLAS_NUM_THREADS=1 in the environment variables.
 - call `openblas_set_num_threads(1)` in the application on runtime.
- You can compile the library itself in sequential mode:
 - make `USE_THREAD=0 USE_LOCKING=1*`(see comment below)
 - If your application is parallelized by OpenMP, please build OpenBLAS with `USE_OPENMP=1`

Next lectures

- Monday 5:
 - Tutorial on STREAM
- Tuesday 6:
 - Tutorial on HPL
- Wednesday 7:
 - Tutorial on IOR/iozone

Ruggero Lot with the help of Niccolo' Tosato.