

Memory allocation

Luca Tornatore - I.N.A.F.



“Foundation of HPC” course



**DATA SCIENCE &
SCIENTIFIC COMPUTING**

2021-2022 @ Università di Trieste

Outline



Memory
Padding



Overhead of
Memory Allocation

Memory allocation has a memory cost

```
typedef struct
{
    char    char_field;
    double  double_field;
    int     int_field;
    char    char_field2;
} STRUCT_A;
```

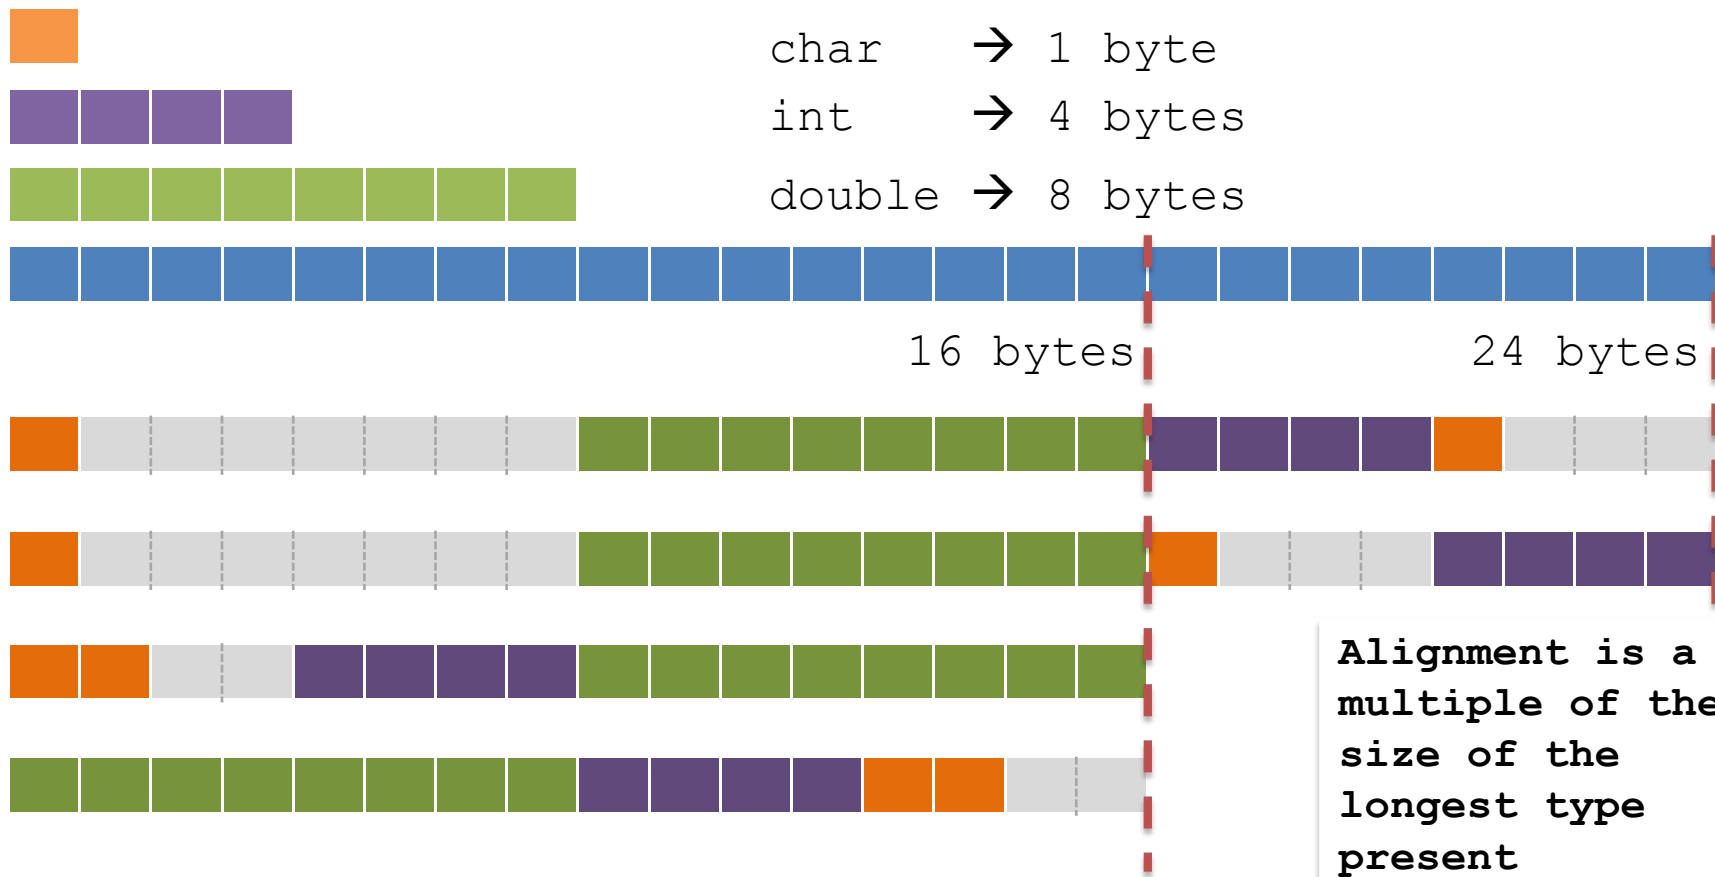
```
typedef struct
{
    char    char_field;
    double  double_field;
    char    char_field2;
    int     int_field;
} STRUCT_B;
```

```
typedef struct
{
    char    char_field;
    char    char_field2;
    int     int_field;
    double  double_field;
} STRUCT_C;
```

```
typedef struct
{
    double  double_field;
    int     int_field;
    char    char_field;
    char    char_field2;
} STRUCT_D;
```

Is there any difference among the above C structures ?

Memory allocation has a memory cost : padding



Memory allocation has a memory cost : padding

```
gcc ... -fpack_struct[=n]
```

An instruction to pack them all and in the bitfield chain them down.

To be used carefully: it generates code *binary-incompatible* with code generated without the option (offset are different) and *sub-optimal code*.

Normally used for non-default binary interface (reduces the data stream).

If given, n must be a (small) power of two.

```
__attribute__ ( (packed) ) ;
```

Inline in the structure definition, at each field you want not to waste any byte.

The same words of caution than above hold.

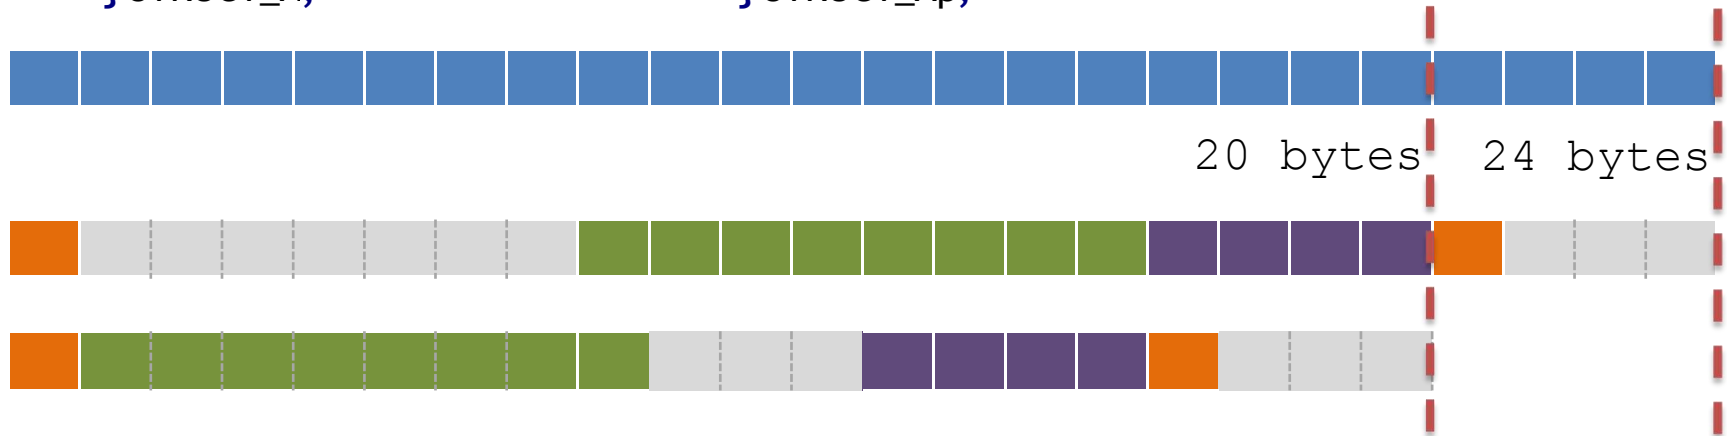
```
typedef struct
{
    char   char_field;
    double double_field;
    int    int_field;
    char   char_field2;
} STRUCT_A;
```

```
typedef struct
{
    char   char_field;
    double double_field __attribute__ ((packed));
    int    int_field;
    char   char_field2;
} STRUCT_Ap;
```

Memory allocation has a memory cost : padding

```
typedef struct
{
    char   char_field;
    double double_field;
    int    int_field;
    char   char_field2;
} STRUCT_A;
```

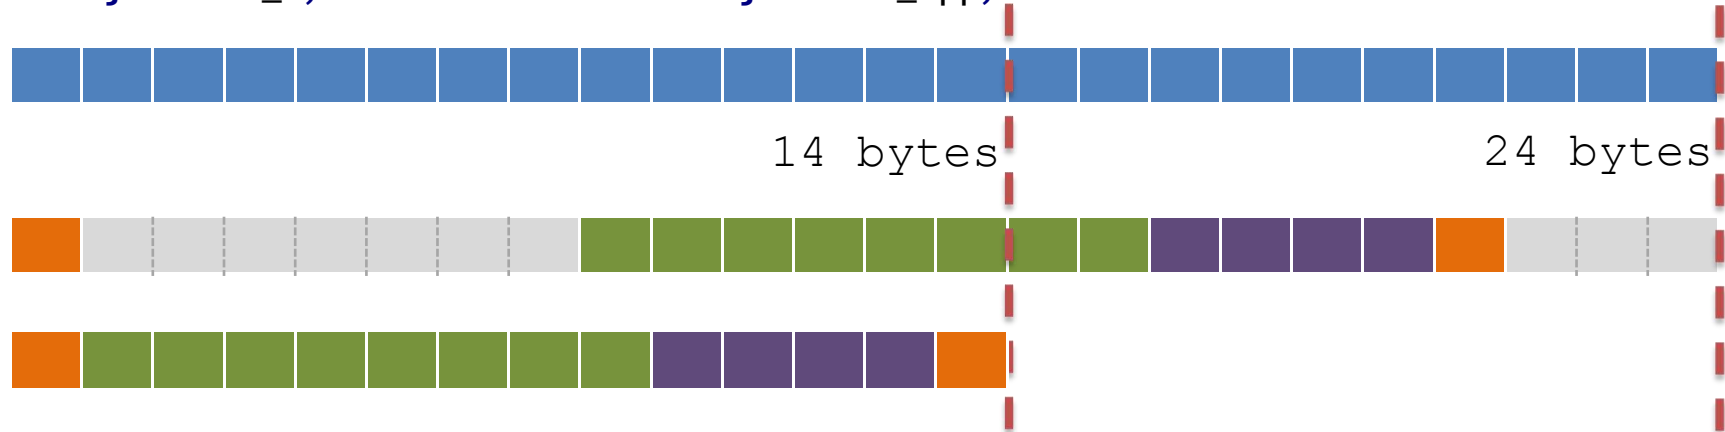
```
typedef struct
{
    char   char_field;
    double double_field __attribute__((packed));
    int    int_field;
    char   char_field2;
} STRUCT_Ap;
```



Memory allocation has a memory cost : padding

```
typedef struct
{
    char   char_field;
    double double_field;
    int     int_field;
    char    char_field2;
} STRUCT_A;
```

```
typedef struct
{
    char   char_field;
    double double_field __attribute__((packed));
    int     int_field __attribute__((packed));
    char    char_field2;
} STRUCT_App;
```



SC0/examples_on_memory_allocation/
structure_padding.c

Memory allocation has a memory cost : padding

General remark:



in order to have **best memory access performances**, it is usually better to have **data aligned** to the natural alignment of your machine, which typically is 64 bits.

```
__attribute__ ((aligned (n))) ;
```

You may achieve this in your data structures by using

Memory allocation has a memory cost : padding

System's `malloc()` allocator returns memory addresses multiples of 8 or 16 for 32- and 64-bits systems.

If you different needs (for instance: using AVXN instructions), you can use special calls:

```
void * posix_memalign (void **memptr, size_t alignment, size_t size)
```

```
void * aligned_alloc (size_t alignment, size_t size)
```

introduced in ISO C11 and hence may have better portability to non-POSIX systems

Outline



Overhead of
Memory Allocation

Memory allocation has a memory cost

Everytime you allocate some memory chunk on the heap, the system has to **keep track** of it, in order to

- maintain a precise account of used and available memory
- be able to “register back” that chunk into the free memory when you don’t need it anymore (you `free()` it).

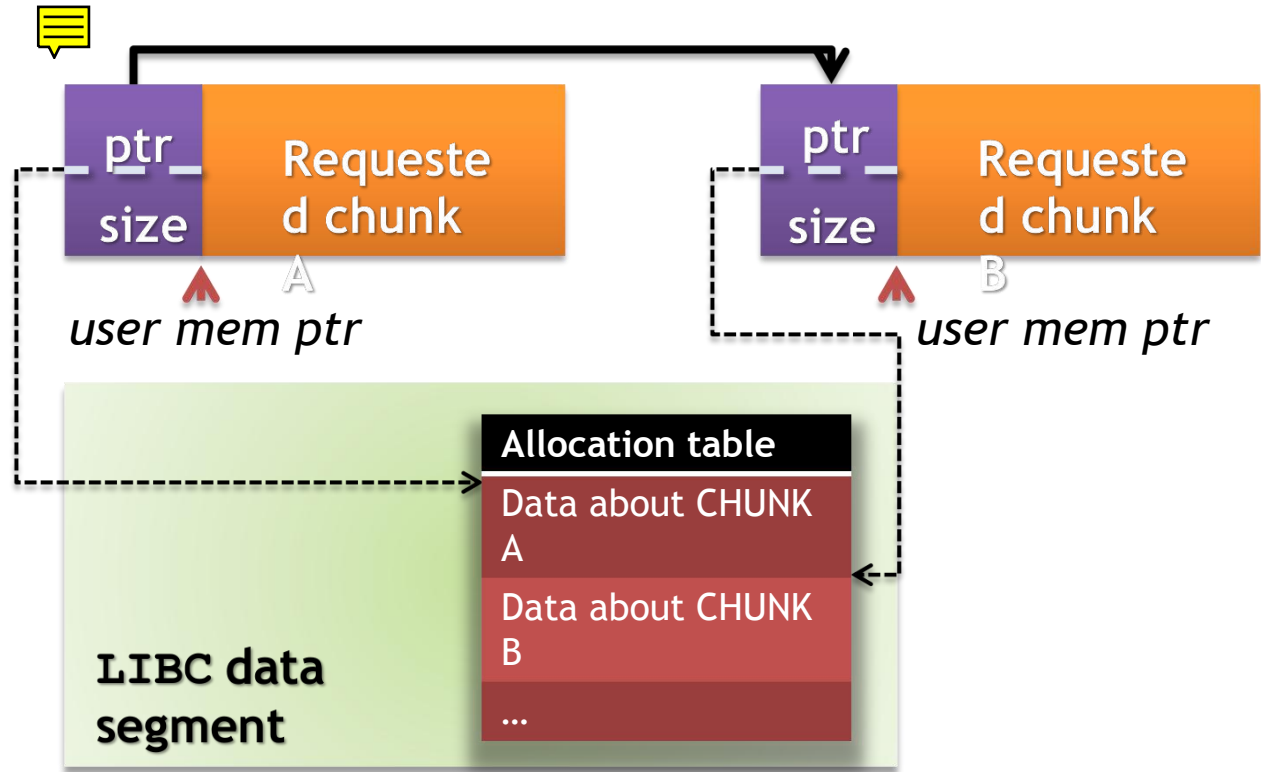
So, in addition to allocate *padded* memory layouts, which amounts to consume more memory than you formally request, there is also an additional “**administrative**” **overhead** whose amount is strongly system- and implementation-dependent (normally few bytes).

Memory allocation has a memory cost

libc prepends an header to the chunk of data it allocates upon your request.

In that header, whose dimension is implementation-dependent, there are several data it uses to check the consistency of memory chain and to recover the data about that same allocation.

Every detail is strongly implementation-dependent !



Memory allocation has a memory cost

Allocating a huge amount of small data (like in text-book linked-list techniques) might be not the optimal strategy:


- you incur in some waste of memory due to padding, depending on the layout of data;
- for sure in a larger amount of data needed to trace every chunk you require.

How much large is the latter overhead ?

Again, that is implementation-dependent.

Let's measure it on your computer and for different implementation
(i.e. *different compilers and libc*)



# of allocated blocks	Requested Mem	Additionally located	Ptr Overhead 	System overhead
10M	400 MB	24 %	80 MB	18 %
1M	400 MB	1.3 %	8 MB	3.2 %
100k	400 MB	0.09 %	0.8 MB	0.27 %
10k	400 MB	0.0 %	80 K	0.05%
1k	400 MB	0.12 %	8 K	0.125 %
100	400 MB	0.033 %	0.8 K	0.034 %
10	400 MB	0.006 %	0.1 K	0.006 %